HAV41@pitt.edu

DBMS - Homework 01

PART-01

Q1 [5 pt]
Find the names of all the instructors from Finance department.

Solution:

Relational Algebra Expression:
Π(name) (σ (dept_name=" finance") (instructor))

SQL Query:
select name from instructor where dept_name='finance';

Output:

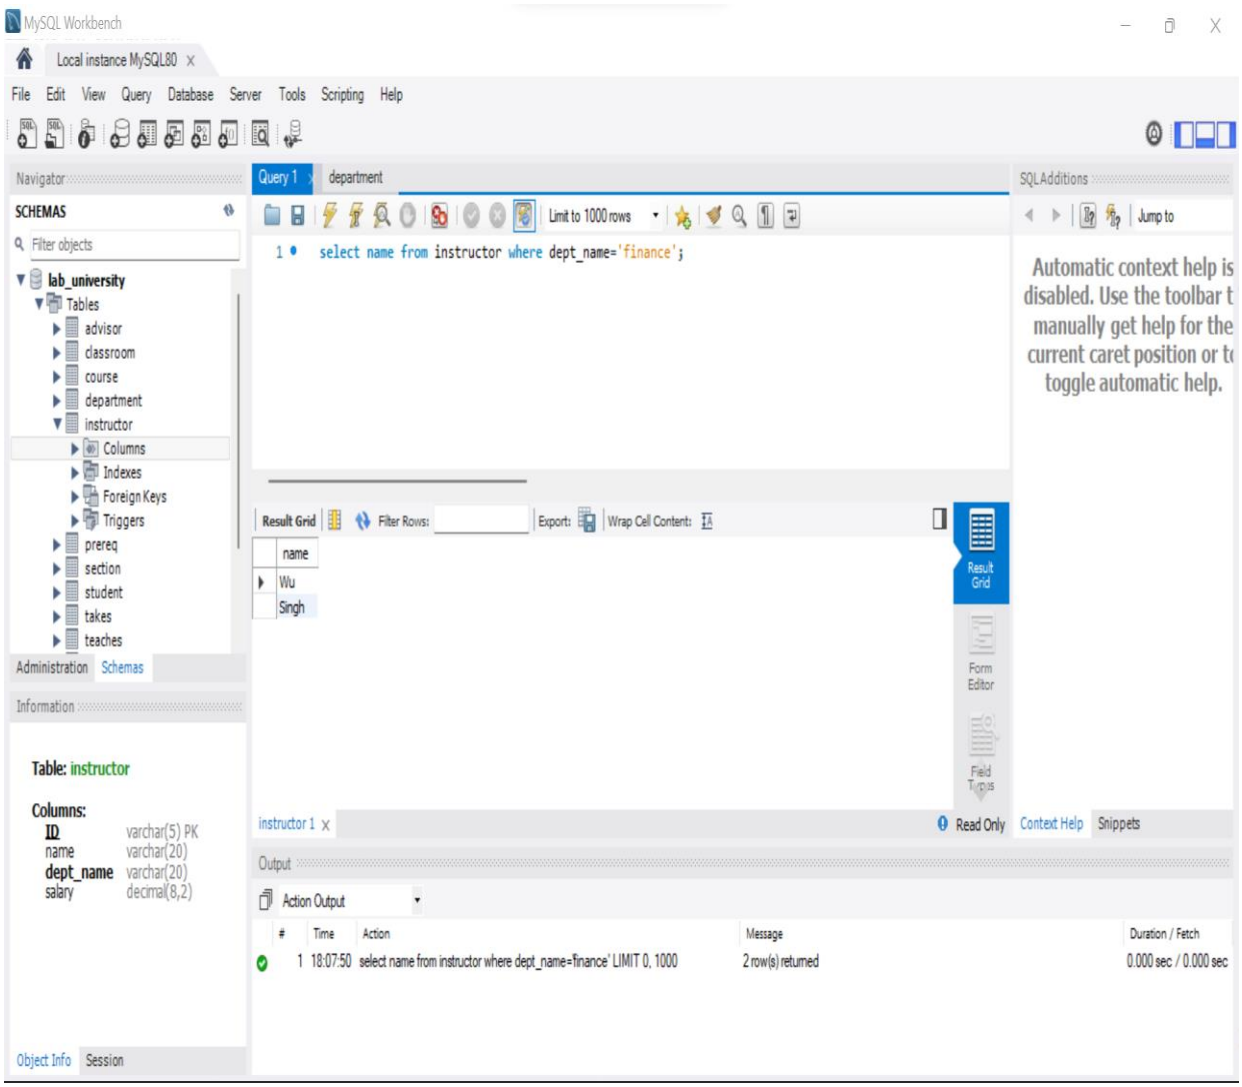

HAV41@pitt.edu

Q2 [5 pt]
Find the names of courses in Biology department which have 4 credits.

Solution:

Relational Algebra Expression:
Π(title)(σ (dept_name=”biology” and credits=”4”)(course))

SQL Query:
select title from course where dept_name='Biology' and credits='4';

Output:

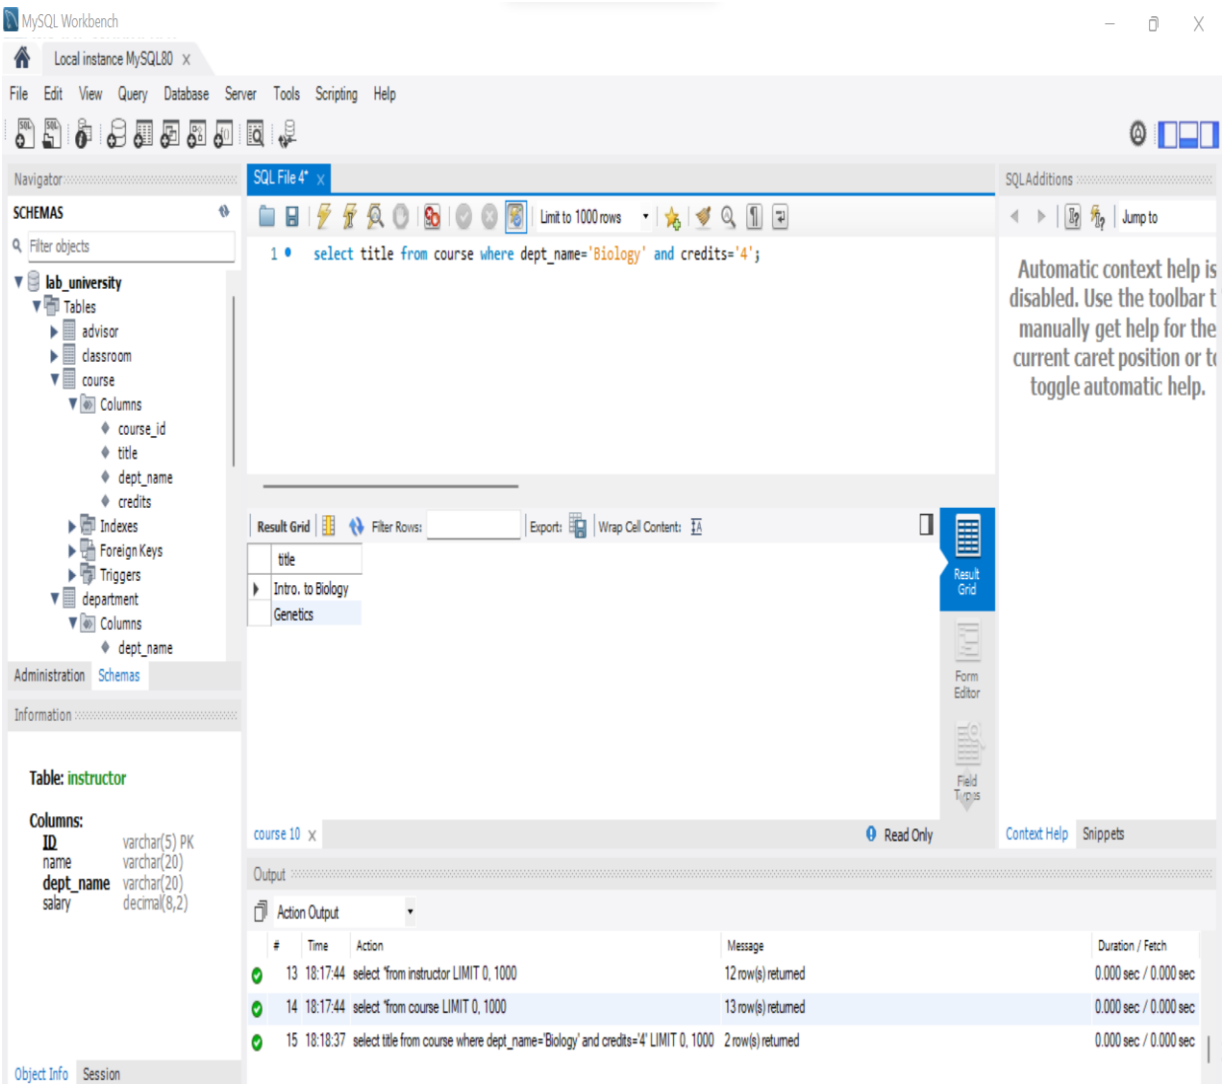

HAV41@pitt.edu

Q3 [5 pt]
Display information about all courses whose name start with letter “I”.

Solution:

Relational Algebra Expression:
Π(course) (σ (title like 'I%')(course))

SQL Query:
select * from course where title like 'I%';

Output:

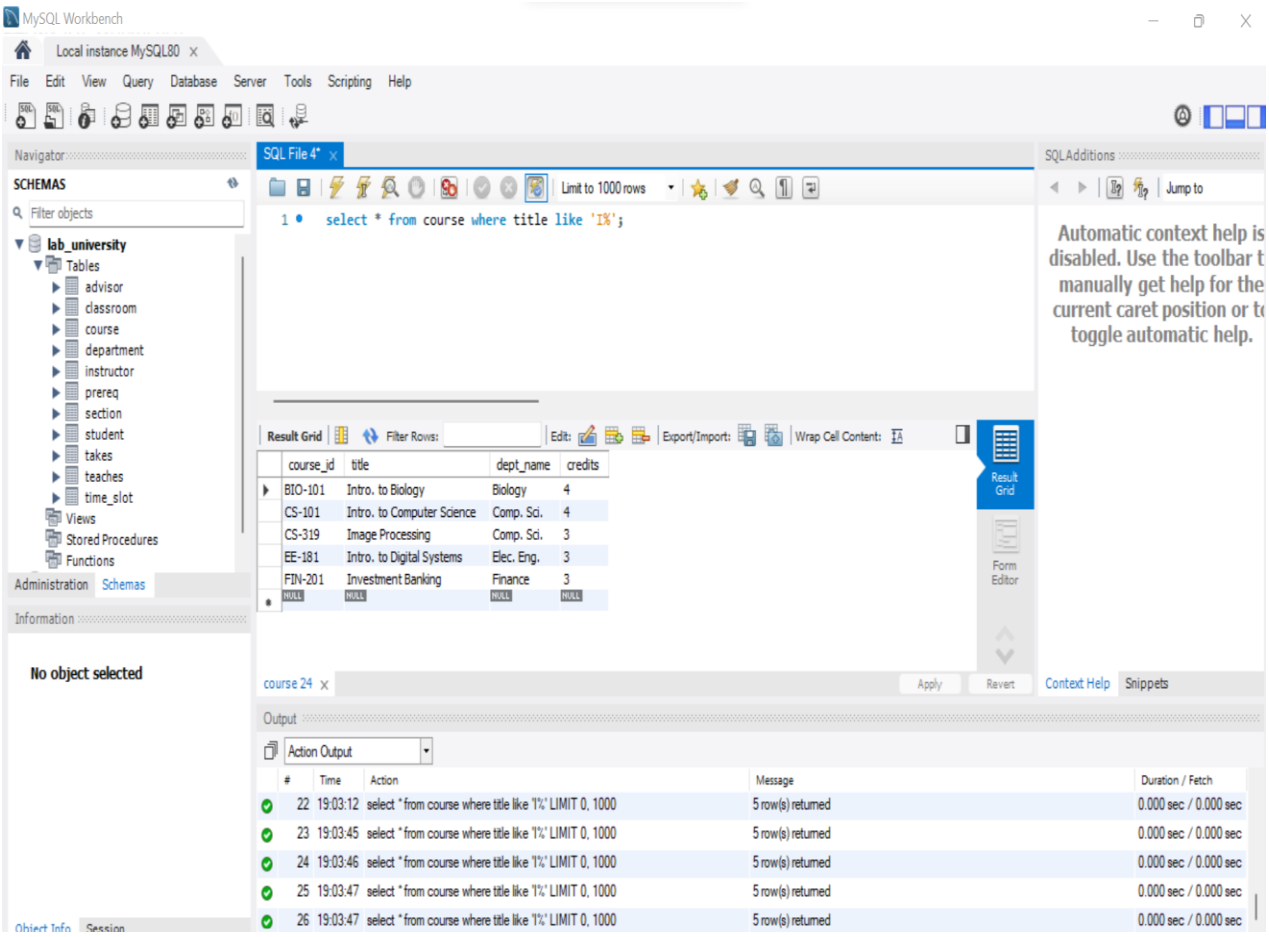

HAV41@pitt.edu

Q4 [5 pt]
Retrieve all course_id and title of all courses taken by the student with ID 12345.

Solution:

Relational Algebra Expression:
Π(takes.course_id,course.title)(σ (ID=’12345’)(course ⋈ takes on course_id))

SQL Query:
select takes.course_id as course_id, course.title as title from course inner join takes on takes.course_id = course.course_id where ID ='12345';

Output:

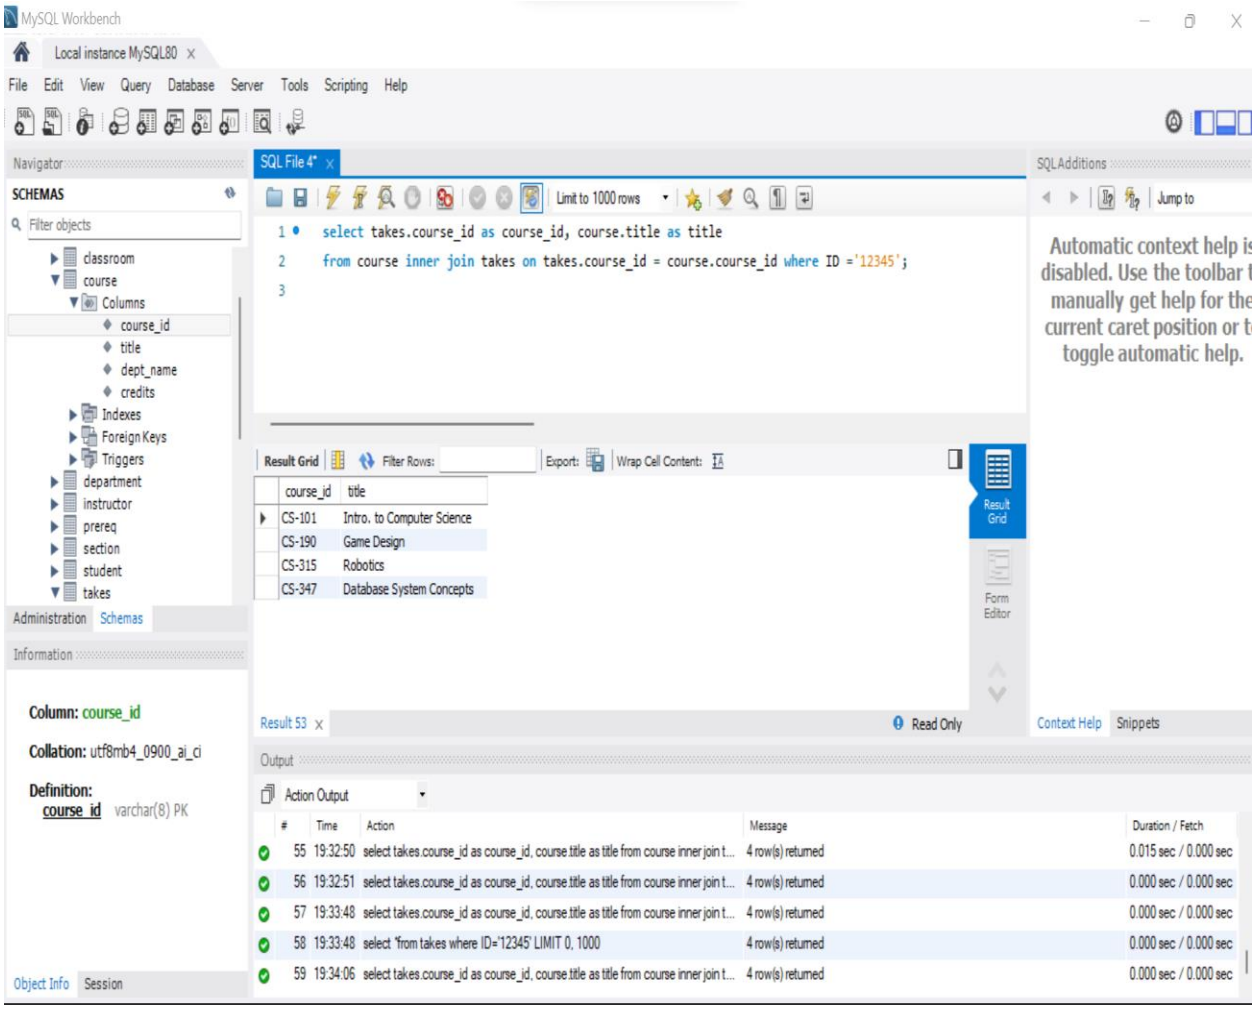

HAV41@pitt.edu

Q5 [10 pt]
Retrieve distinct course_id and names of all the courses taught in 2009 Spring and 2010 Spring.

Solution:

Relational Algebra Expression:
Π(section.course_id,course.title)(σ (year in (2009,2010)) and semester = “Spring”)(course ⋈ section on course_id))

SQL Query:

```
select distinct section.course_id, course.title as title from course inner join
section on course.course_id = section.course_id where year in (2009,2010) and semester = 'Spring';
```

Output:

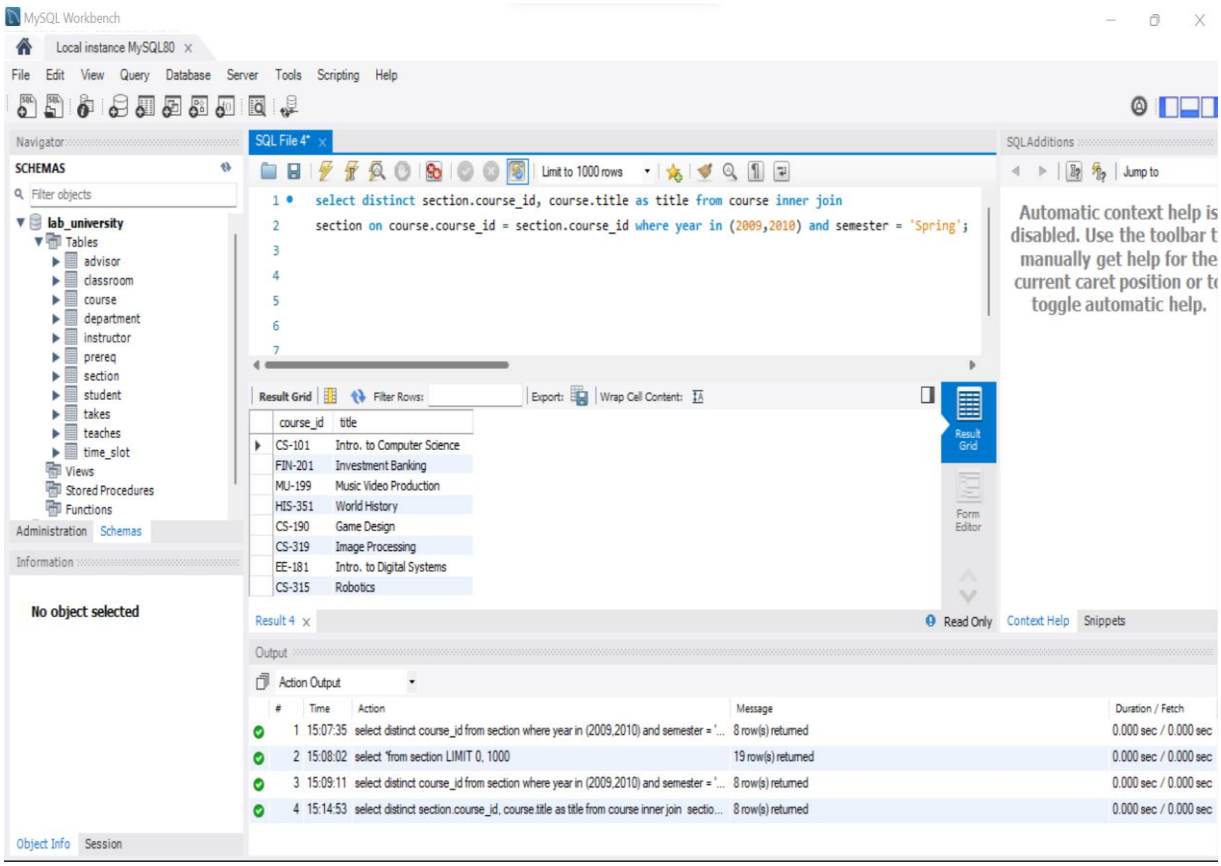

HAV41@pitt.edu

Q6 [10 pt]
Find course_id and name of all courses that taught by instructor from history department.

Solution:

Relational Algebra Expression:
Π(course.course_id,course.title)(σ dept_name = “History”)(course ⋈ instructor on dept_name) ⋈ teaches on ID)

SQL Query:

```
select course.course_id, title from course inner join instructor on course.dept_name =
instructor.dept_name  inner join teaches on instructor.ID = teaches.ID where
instructor.dept_name = 'History';
```

Output:

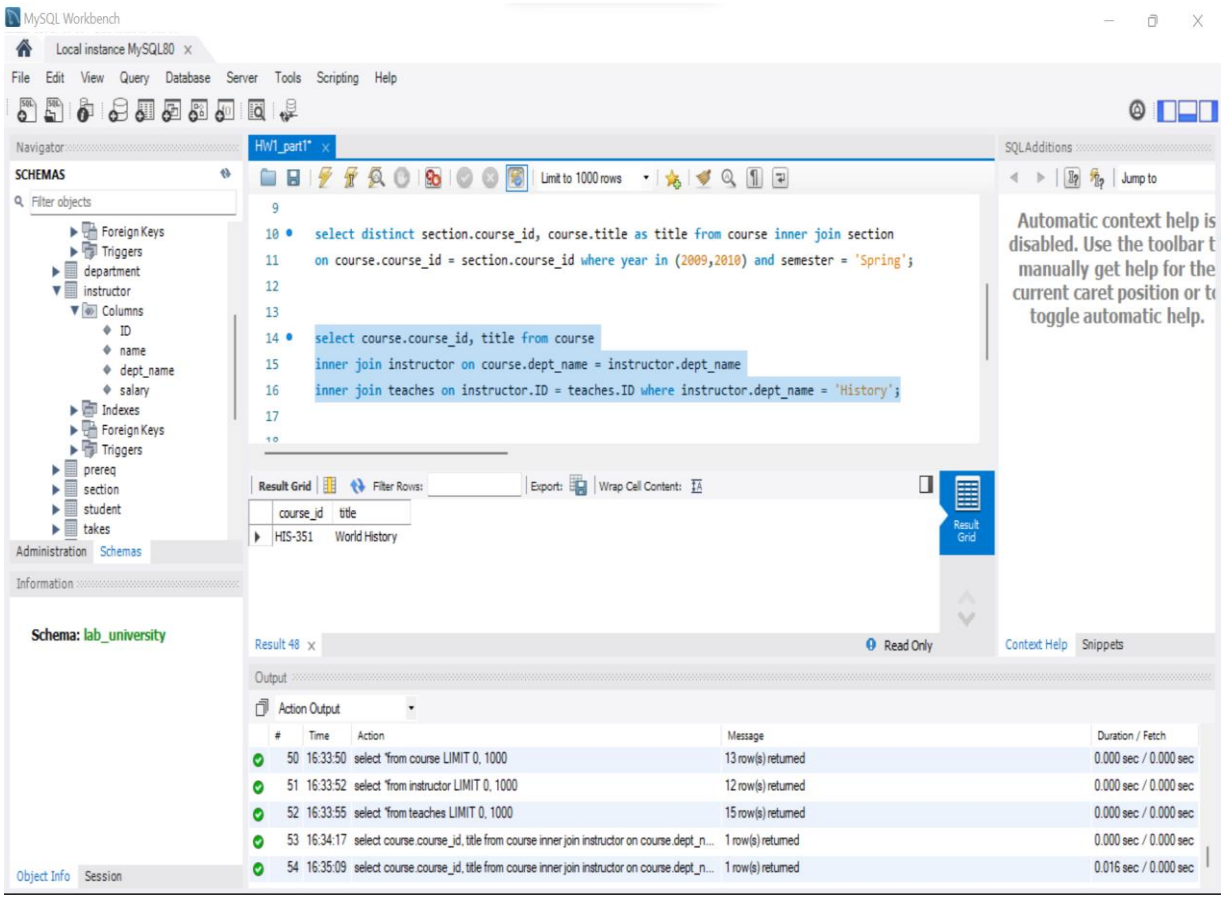

HAV41@pitt.edu

PART-02

Q7[15 pt]
Write SQL DDL statements to create the above tables. Make sure that you capture the primary and foreign key constraints (if applicable), choose appropriate domain (data) type and constraints for each attribute.

Solution:

Salesman Table

SQL Query to create table:

```
create table salesman (salesman_id int,
                       salesman_name varchar(30),
                       city     varchar(30),
                       commission    decimal(10, 2),
                       primary key (salesman_id));
```

SQL Query to insert values into the table:

```
insert into salesman values ('5001', 'James Hoog', 'New York','0.15');
insert into salesman values ('5002', 'Nail Knite', 'Paris','0.13');
insert into salesman values ('5005', 'Pit Alex', 'London','0.11');
insert into salesman values ('5006', 'Mc Lyon', 'Paris','0.14');
insert into salesman values ('5003', 'Lauson Hen', null,'0.12');
insert into salesman values ('5007', 'Paul Adam', 'Rome','0.13');
```

Output:

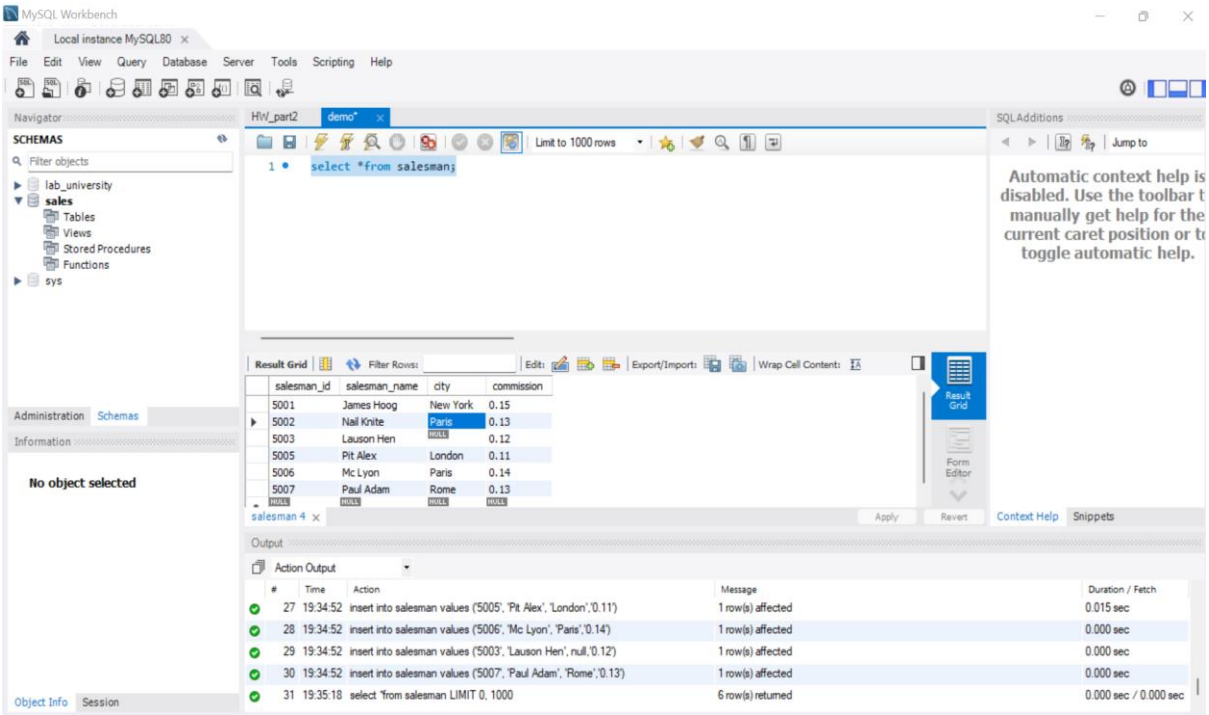

HAV41@pitt.edu

Customer Table

SQL Query to create table:

```
create table customer (customer_id int,
                       customer_name varchar(30),
                       city     varchar(30),
                       grade  int,
                      primary key (customer_id));
```

SQL Query to insert values into the table:

```
insert into customer values ('3002', 'Nick Rimando', 'New York','100');
insert into customer values ('3005', 'Brad Guzan', 'London','200');
insert into customer values ('3001', 'Brad Davis', 'New York',null);
insert into customer values ('3004', 'Fabian Johns', 'Paris','300');
```

Output:

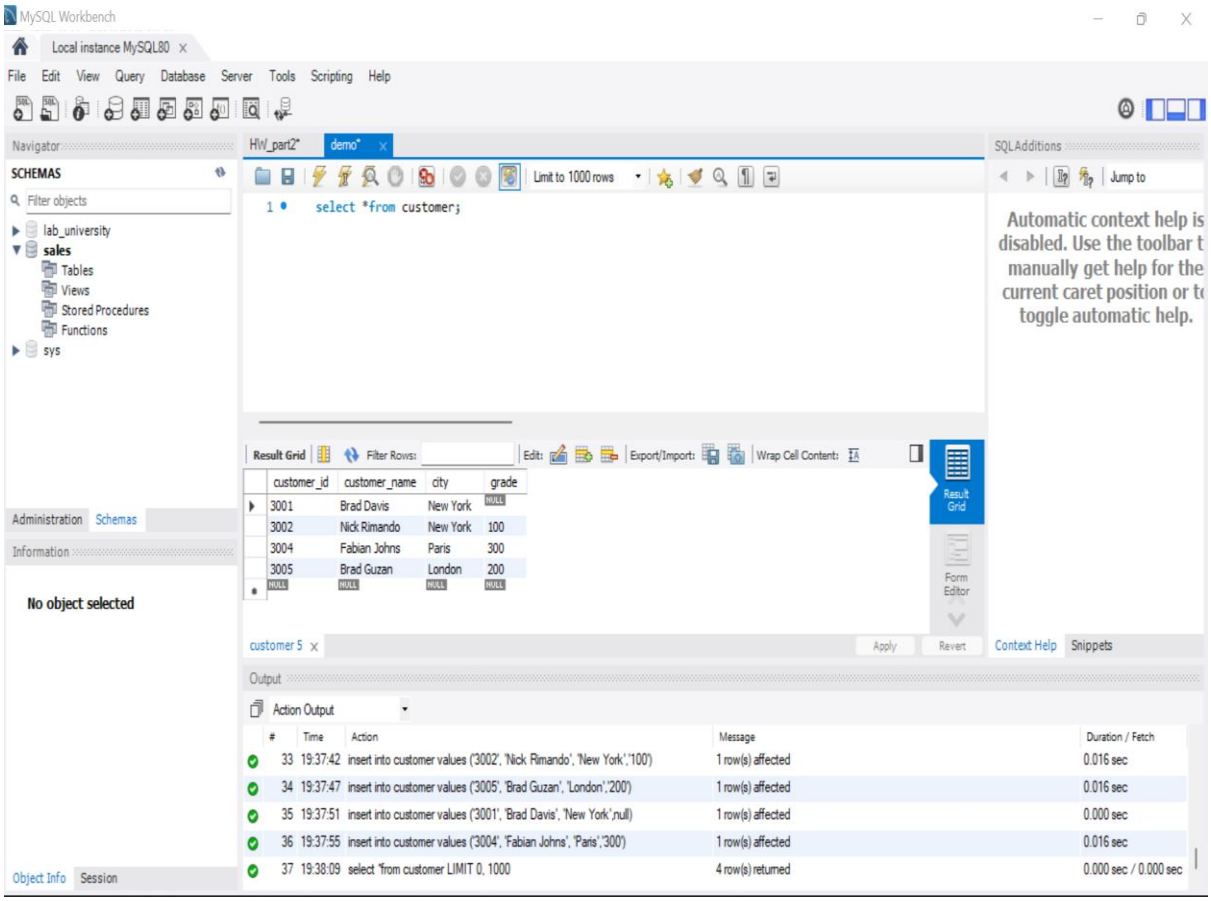

HAV41@pitt.edu

Order Table

SQL Query to create table:

```
create table orders (order_noint,
                     purchase_amount decimal(10, 2),
                     order_date date,
                     customer_id int,
                     salesman_id int,
                     primary key (order_no),
          foreign key (customer_id) references customer(customer_id) on delete
cascade,
          foreign key (salesman_id) references salesman(salesman_id) on delete cascade
        );
```

SQL Query to insert values into the table:

```
insert into orders values ('7001', '150.5', '2016-06-27','3005','5002');
insert into orders values ('7002', '270.65', '2016-07-17','3002','5001');
insert into orders values ('7003', '110.5', '2016-07-27','3005','5002');
insert into orders values ('7004', '948.5', '2016-10-05','3004','5006');
insert into orders values ('7005', '250.45', '2016-10-10','3001',null);
insert into orders values ('7006', '79.29', '2016-10-10','3002','5001');
```

Output:

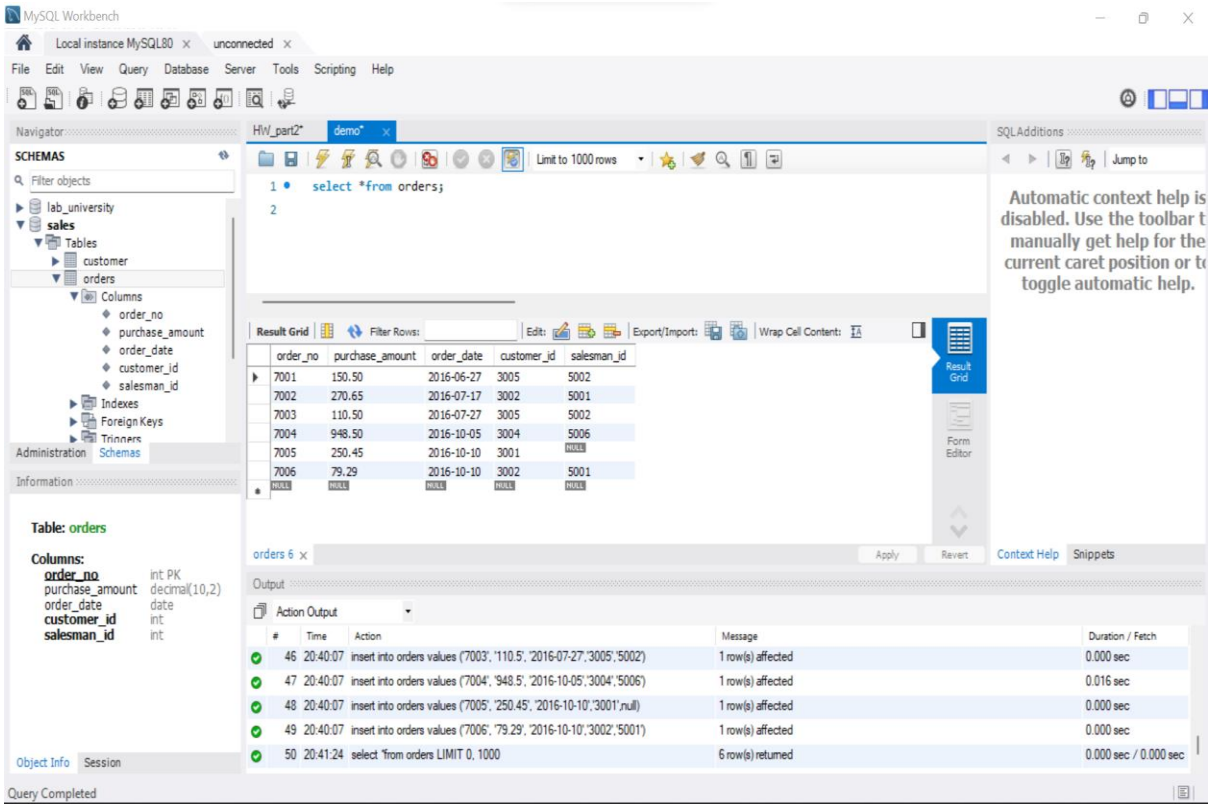

HAV41@pitt.edu

Q8 [5 pt]
Retrieve customer id of all customers from orders table without any repeats.

Solution:

Relational Algebra Expression:
Π(customer_id)(σ (orders))

SQL Query:
select distinct customer_id from orders;

Output:

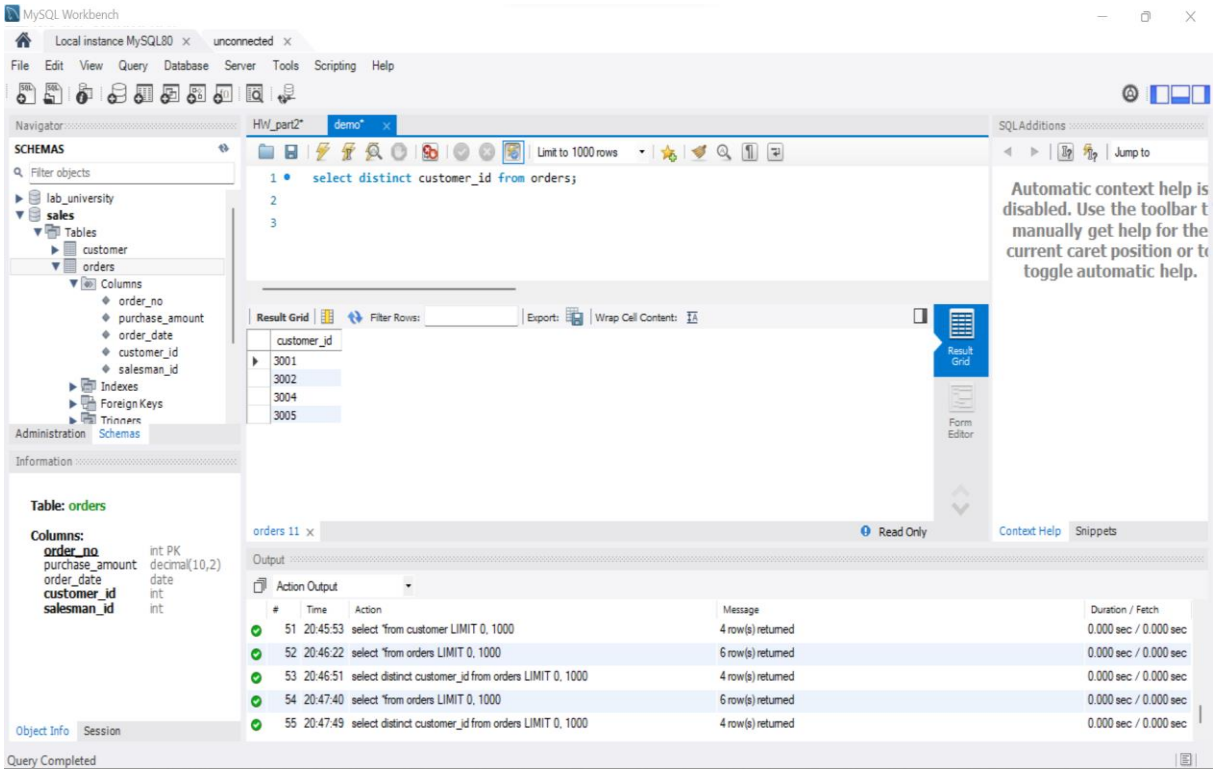

HAV41@pitt.edu

Q9 [5 pt]
Display the names and city of salesman, who doesn't belong to the city of Paris.

Solution:

Relational Algebra Expression:
Π(customer_name,city)(σ (city≠Paris)(customer))

SQL Query:
select customer_name, city from customer where not city = 'Paris';

Output:

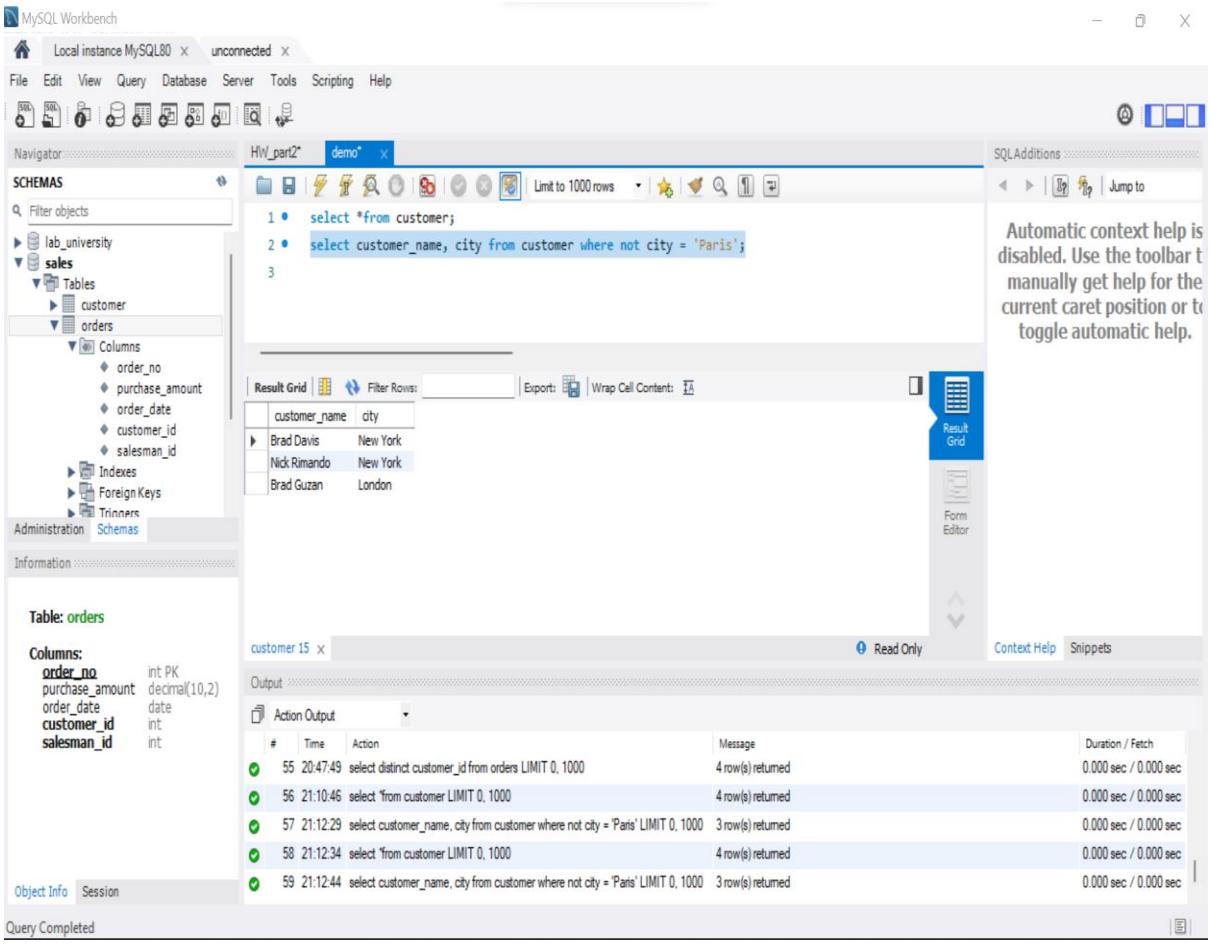

HAV41@pitt.edu

Q10 [5 pt] Find that customer with all information who does not get any grade except NULL.

Solution:

Relational Algebra Expression:
Π(customer_id,customer_name,city)(σ (grade=Null)(customer))

SQL Query:
select customer_id, customer_name,city from customer where grade is null;

Output:

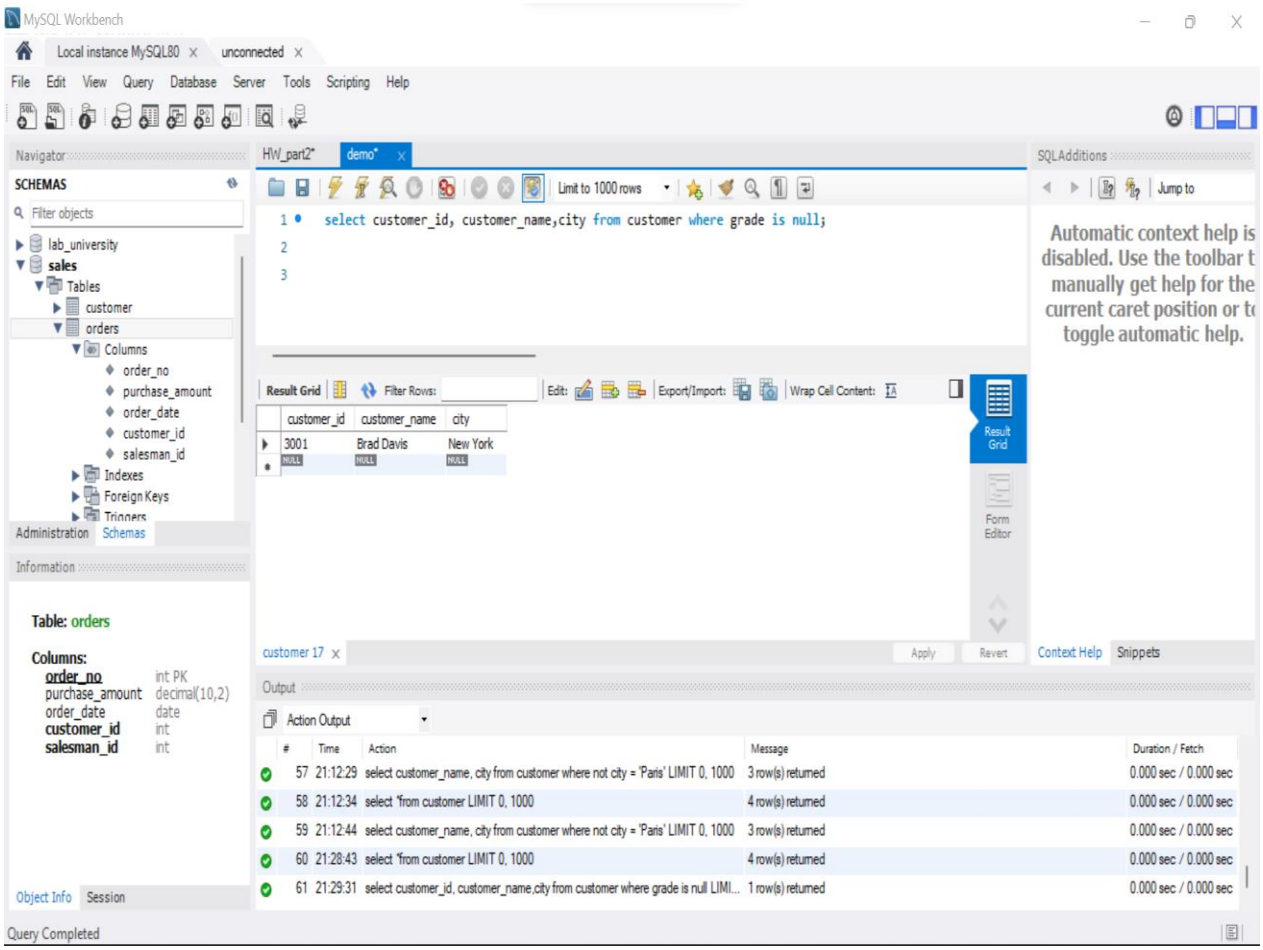

HAV41@pitt.edu

Q11[20 pt] Consider the following relational algebra expression.

$$\pi_{salesman\_id,name,city,commission}(\sigma_{(commission \geq 0.10) \cap (commission \leq 0.12)} salesman)$$

1) How many attributes will the result have?

   Solution: Salesman_id , name , city and commission → 4 attributes.

2) Write in English what question the expression is trying to answer (e.g. describe what would be the result of the expression).

   Solution:
   Select salesman id, name, city, and commission from salesman table where commission is greater than or equal to 0.10 and commission is lesser than or equal to 0.12. The output will have four columns that is salesman id, name, city and commission. The row data will be based on the condition → commission is greater than or equal to 0.10 and commission is lesser than or equal to 0.12. In simpler terms salesman id, name and city of the individual will be displayed based on the above condition.

3) Translate the expressions into SQL.
   Solution:

   ```
   select salesman_id,name,city,commission from
   salesman where commission >=0.10 and commission<=0.12
   ```